2 | Normalised attention entropy by layer

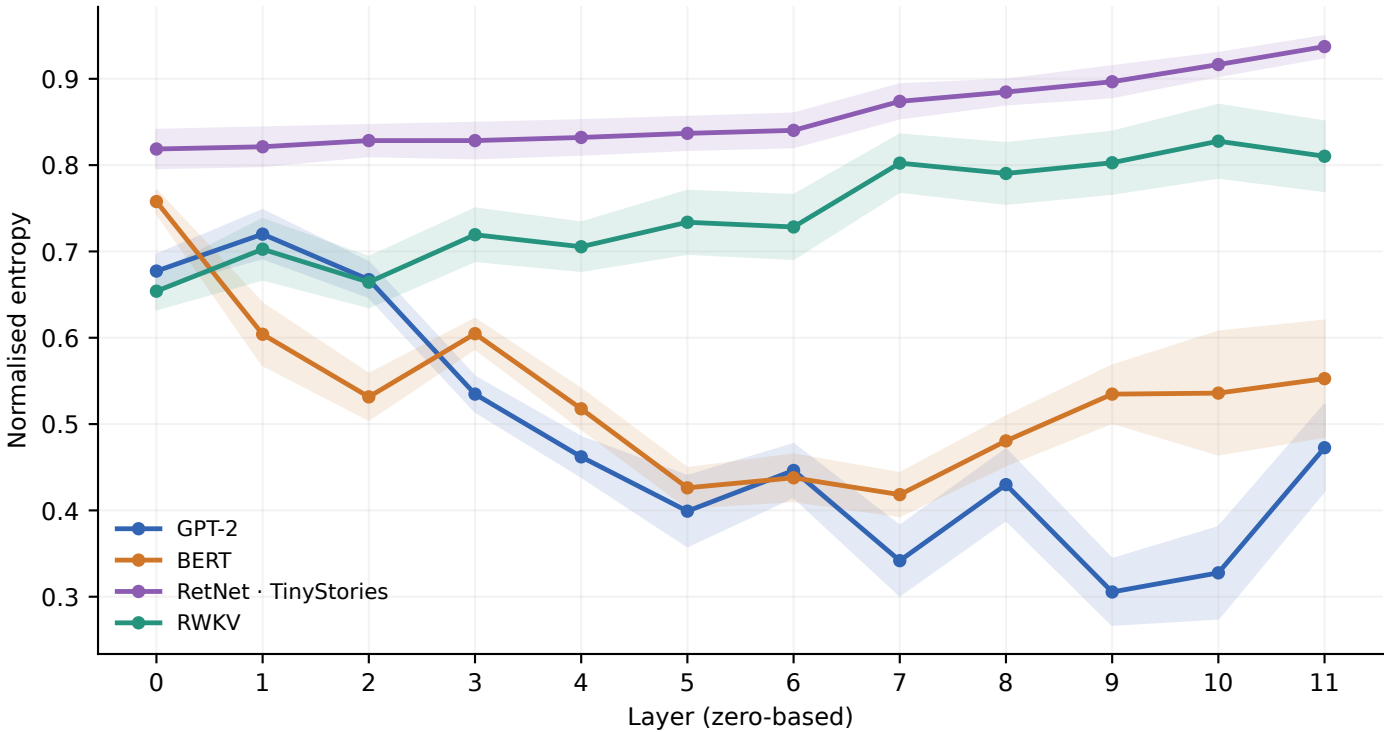

Final-layer summary

Model	Mean	SD
GPT-2	0.47	0.05
BERT	0.55	0.07
RetNet · TinyStories	0.94	0.01
RWKV	0.81	0.04

32 identical prompts per model; lines = prompt mean; shaded bands = ±1 population SD, not confidence intervals.
Entropy excludes the first query for causal models.