5 | First-token attention mass by layer

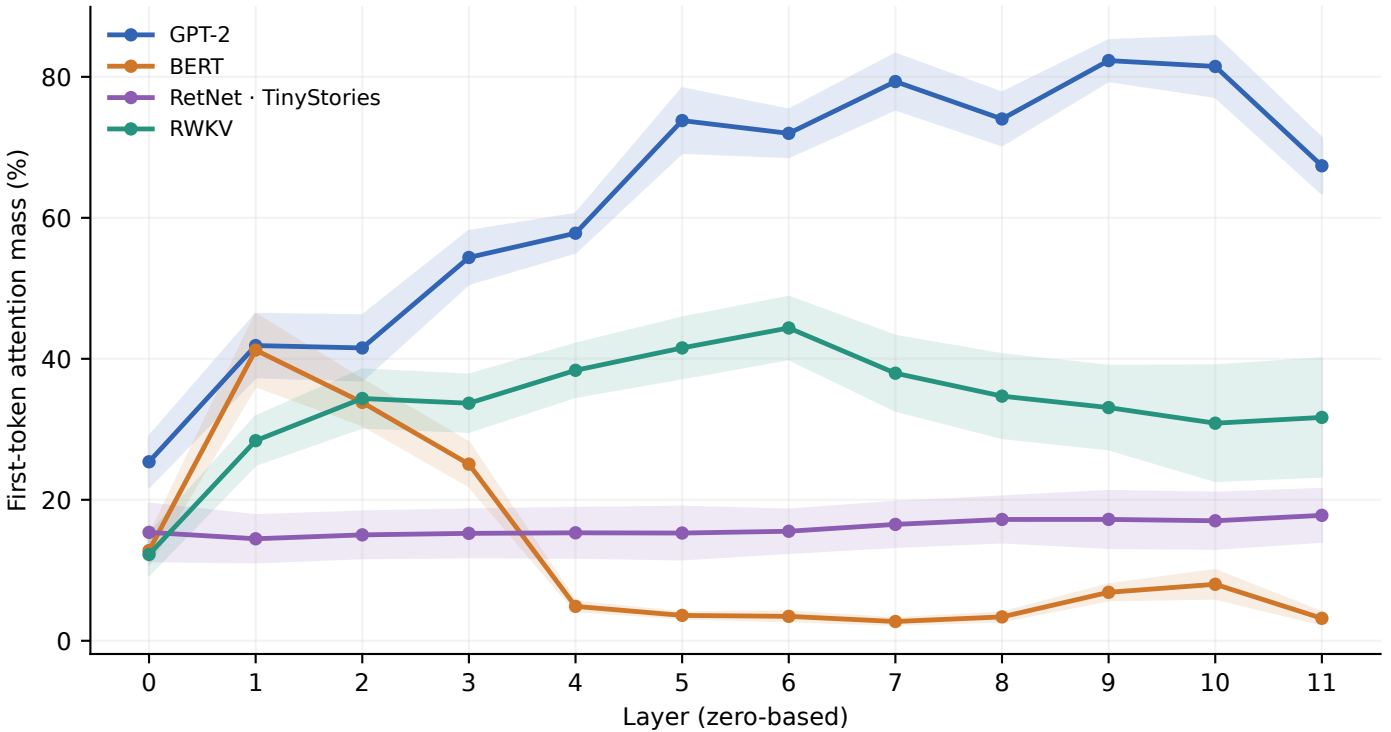

Final-layer summary

Model	Mean	SD
GPT-2	67.38	3.89
BERT	3.19	0.83
RetNet · TinyStories	17.79	3.71
RWKV	31.69	8.35

32 identical prompts per model; lines = prompt mean; shaded bands = ±1 population SD, not confidence intervals.
First query excluded; BERT's first token is [CLS], while causal models use the first text token.